

Transfer Learning

Transfer Learning is a deep learning technique that reuses knowledge from a source task to enhance learning on a new, related target task. It addresses the significant challenges of training deep learning models from scratch, which typically demands:

Huge Datasets

Requiring millions of labeled samples.

High Compute

Demanding expensive GPUs/TPUs.

Extended Time

Often taking days to weeks.

Specialized Expertise

Needing careful model tuning.

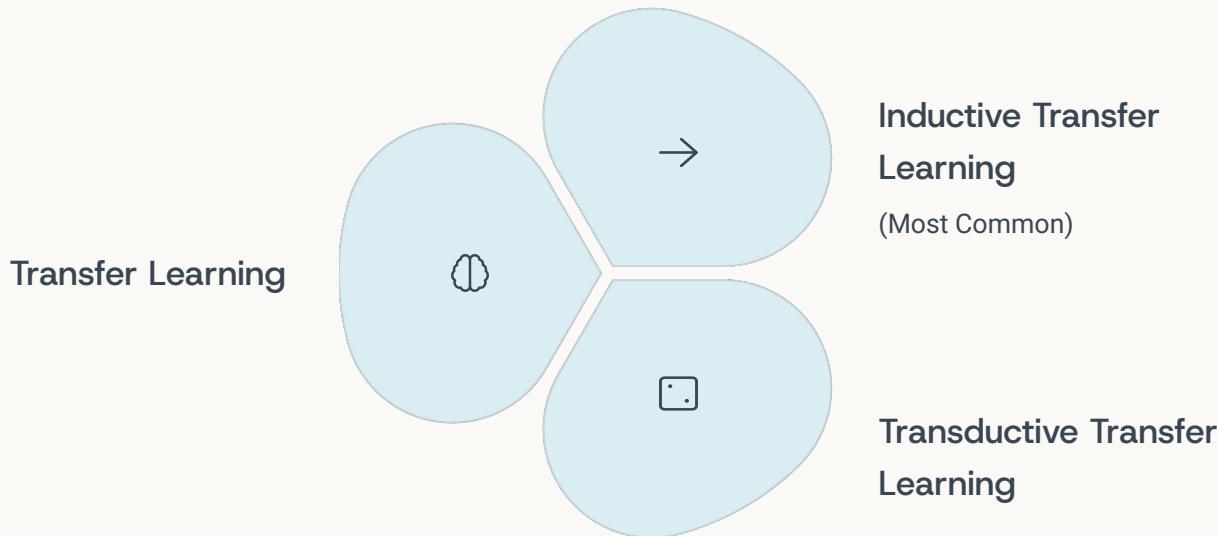
Transfer Learning is your go-to strategy when:

- **Labeled Data is Limited:** This is the most common scenario.
- **A Related Pretrained Model Exists:** Leveraging existing knowledge is key.
- **Training from Scratch is Costly:** Reducing computational burden.
- **Task Similarity:** The new task shares characteristics with a known task.



Transfer Learning: Key Types

Transfer Learning encompasses various approaches, each suited for different scenarios. Understanding these types helps in selecting the most effective strategy for your specific problem.



These categories define how source and target tasks, domains, and data labels influence the transfer process. We'll delve into the most prevalent types to illustrate their applications.

Inductive Transfer Learning

Inductive Transfer Learning is applied when the source and target tasks are different, but crucially, labeled data is available for the target task. Both source and target tasks possess distinct output spaces, objectives, and loss functions.

1

Source Task

Dataset: ImageNet

Task: Classify 1000 object categories

Model: ResNet50 (pretrained)

2

Target Task

Dataset: Plant leaf images

Task: Detect plant diseases (Healthy / Diseased)

Labels: Available

Here's how you'd typically implement Inductive Transfer Learning:

- Load the ResNet50 model, pretrained on ImageNet.
- Remove the original 1000-class classification layer.
- Add a new, custom classifier tailored for plant disease detection.
- Train this new head (using feature extraction or fine-tuning) on your plant leaf dataset.

The key distinction: the tasks are different, but target labels exist, making it a prime candidate for Inductive Transfer Learning.

Transductive Transfer Learning

Transductive Transfer Learning, often called domain adaptation, is employed when the source and target tasks are the same, but the data domains differ significantly, and critically, target labels are not available. It's about taking a known task and applying it to new, visually different data distributions.

1

Source Task

Data: English movie reviews

Task: Sentiment analysis

Labels: Yes, sentiment labels are available

2

Target Task

Data: Hindi movie reviews

Task: Sentiment analysis (same task)

Labels: No, sentiment labels are not available

In this scenario, the task (sentiment analysis) remains constant, but the data distribution (language) changes, and there are no labels for the target domain. This perfectly illustrates Transductive Transfer Learning, where the goal is to adapt the model to the new data distribution without direct supervision on the target domain.

Transfer Learning vs. Fine-Tuning: A Closer Look

It's common to conflate Transfer Learning and Fine-Tuning, but it's important to understand their hierarchical relationship: fine-tuning is a specific technique utilized within the broader framework of transfer learning.

Fine-tuning is a technique used ***within*** transfer learning.

Transfer learning is the broader concept; fine-tuning is one way to implement it.



Transfer Learning

The overarching concept of reusing pre-trained models.



Feature Extraction

A method within Transfer Learning where the pre-trained model's layers are frozen.



Fine-Tuning

A method where the pre-trained model's later layers are unfrozen and trained.

Both feature extraction and fine-tuning are strategies for leveraging the "knowledge" (learned weights) from a pre-trained model to accelerate and improve the learning process on a new task.

The Transfer Learning Journey: A Step-by-Step Guide

Implementing Transfer Learning involves a systematic process to effectively adapt a pre-trained model to your specific needs.

Decide if TL is Appropriate

Assess if your problem meets the criteria for effective transfer learning (limited data, related pre-trained model, etc.).

Load Pretrained Weights

Initialize your chosen model with the weights from its source task.

Add New Task-Specific Head

Append a new set of layers tailored to your target task's output requirements.

Compile the Model

Select an optimizer, an appropriate loss function, and a learning rate suitable for your task.

Evaluate and Adjust

Assess the model's performance and make necessary adjustments to hyperparameters or architecture.

Choose a Pretrained Model

Select a robust model trained on a large, relevant dataset (e.g., ResNet, MobileNet, EfficientNet).

Remove Original Head

Discard the final classification or output layer of the pre-trained model, as it's task-specific.

Choose Transfer Strategy

- **Feature Extraction:** Freeze all base model layers, training only the new head.
- **Fine-Tuning:** Unfreeze and train a few of the last layers of the base model along with the new head.

Train on Target Data

Train the adapted model using your specific target dataset.

Popular Pretrained Models

- VGG16 / VGG19

Learning, baselines

- ResNet

General-purpose, stable

- InceptionV3

Multi-scale features

- Xception

Efficient deep CNN

- MobileNetV1 / V2 /

Mobile, Lightweight models

- SqueezeNet

Tiny models

Thank You